

Machine Learning – Lecture 19

Recurrent Neural Networks

09.01.2020

Bastian Leibe


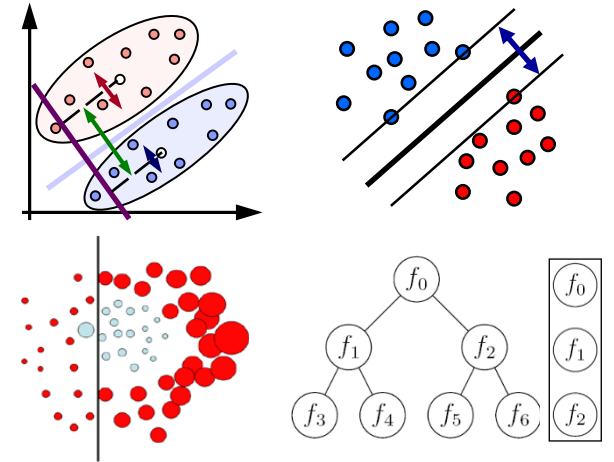

RWTH Aachen

<http://www.vision.rwth-aachen.de>

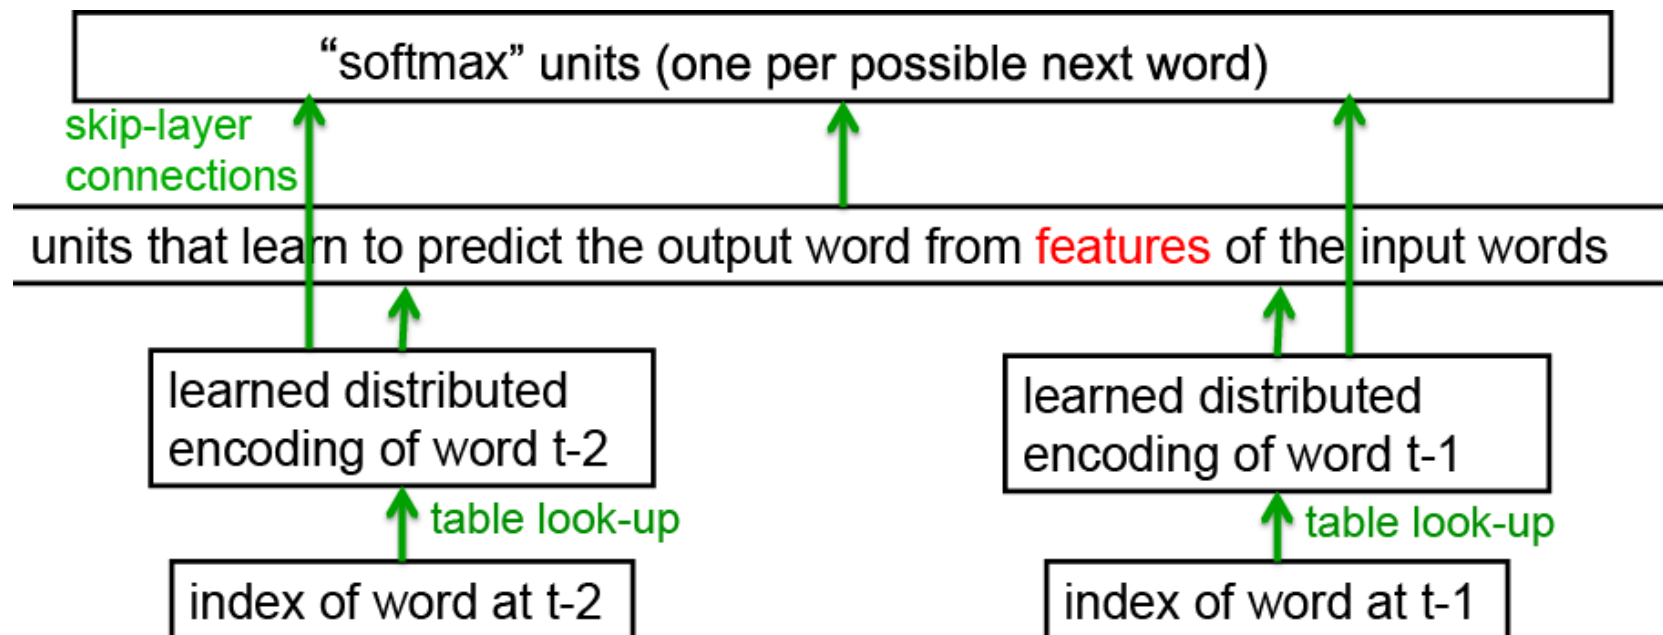
leibe@vision.rwth-aachen.de

Course Outline

- Fundamentals
 - Bayes Decision Theory
 - Probability Density Estimation
- Classification Approaches
 - Linear Discriminants
 - Support Vector Machines
 - Ensemble Methods & Boosting
 - Random Forests
- Deep Learning
 - Foundations
 - Convolutional Neural Networks
 - Recurrent Neural Networks



Recap: Neural Probabilistic Language Model

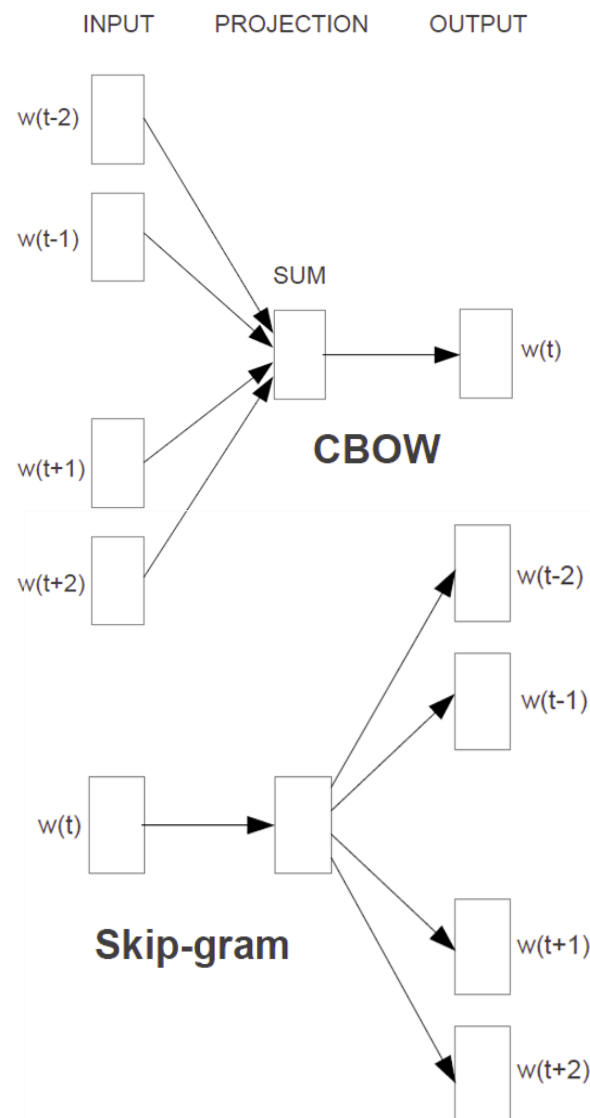


- Core idea
 - Learn a shared distributed encoding (word embedding) for the words in the vocabulary.

Y. Bengio, R. Ducharme, P. Vincent, C. Jauvin, [A Neural Probabilistic Language Model](#), In JMLR, Vol. 3, pp. 1137-1155, 2003.

Recap: word2vec

- Goal
 - Make it possible to learn high-quality word embeddings from huge data sets (billions of words in training set).
- Approach
 - Define two alternative learning tasks for learning the embedding:
 - “Continuous Bag of Words” (CBOW)
 - “Skip-gram”
 - Designed to require fewer parameters.




Recap: word2vec CBOW Model

- Continuous BOW Model


- Remove the non-linearity from the hidden layer
- Share the projection layer for all words (their vectors are averaged)

⇒ Bag-of-Words model
(order of the words does not matter anymore)



Recap: word2vec Skip-Gram Model

- Continuous Skip-Gram Model
 - Similar structure to CBOW
 - Instead of predicting the current word, predict words within a certain range of the current word.
 - Give less weight to the more distant words



Problems with 100k-1M outputs

- Weight matrix gets huge!
 - Example: CBOW model
 - One-hot encoding for inputs
 - ⇒ Input-hidden connections are just vector lookups.
 - This is not the case for the hidden-output connections!
 - State h is not one-hot, and vocabulary size is 1M.
 - ⇒ $\mathbf{W}'_{N \times V}$ has $300 \times 1\text{M}$ entries
- Softmax gets expensive!
 - Need to compute normalization over 100k-1M outputs


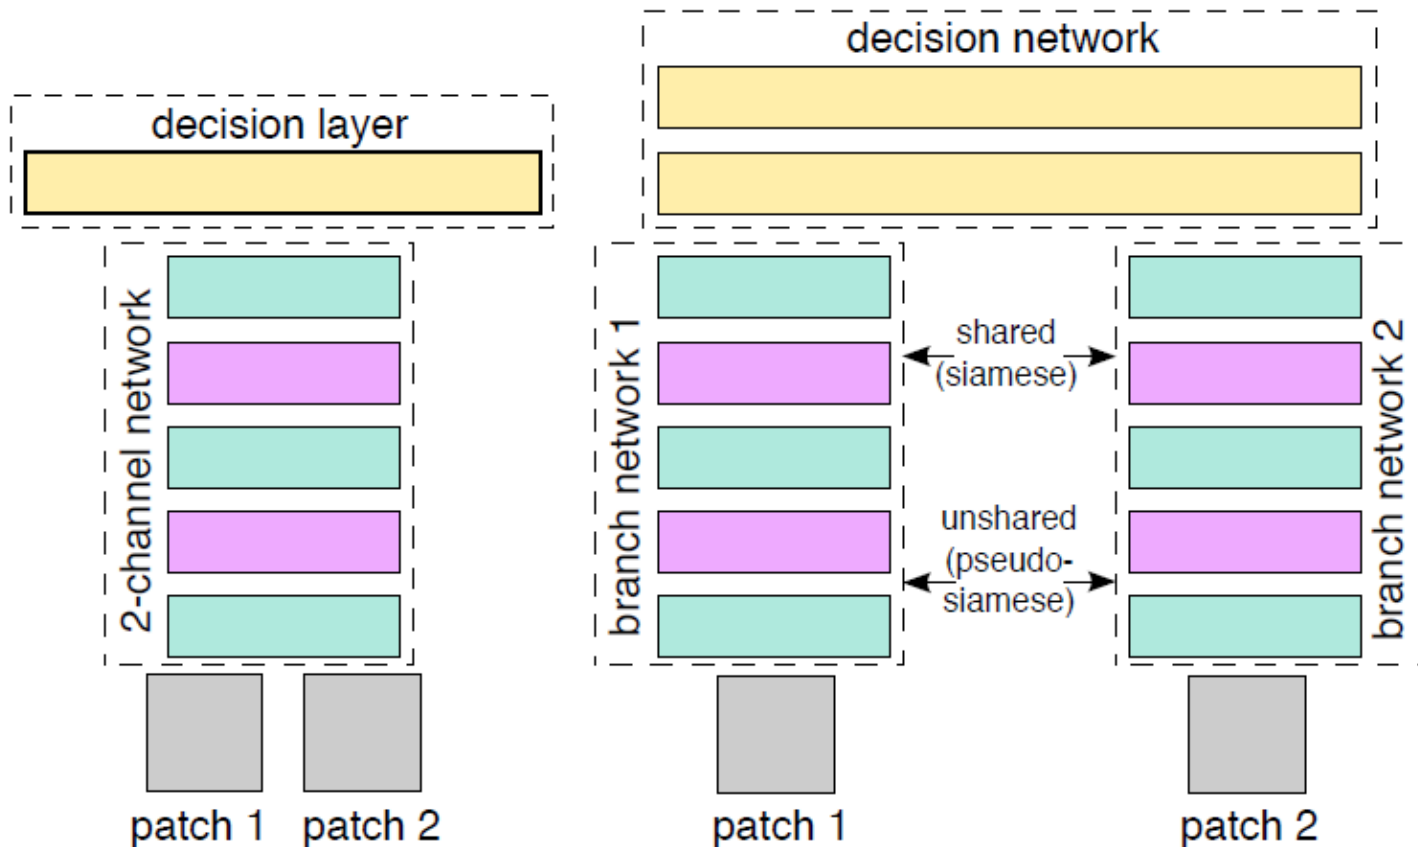


Diagram illustrating a binary tree structure representing a sequence of nodes. The root node is labeled $n(w_2, 1)$. Its left child is labeled $n(w_2, 2)$, and its right child is unlabeled. The node $n(w_2, 2)$ has a left child labeled $n(w_2, 3)$ and a right child labeled w_3 . The node $n(w_2, 3)$ has a left child labeled w_1 and a right child labeled w_2 . The node w_2 is highlighted with a thick black border. The tree continues with more nodes, indicated by an ellipsis between w_4 and w_{V-1} .

- Organize words in binary search tree, words are at leaves
- Factorize probability of word w_0 as a product of node probabilities along the path.
- Learn a linear decision function $y = v_{n(w,j)} \cdot h$ at each node to decide whether to proceed with left or right child node.

⇒ Decision based on output vector of hidden units directly.

Siamese Networks




- Similar idea to word embeddings
 - Learn an embedding network that preserves (semantic) similarity between inputs
 - E.g., used for patch matching

Topics of This Lecture

- Recurrent Neural Networks (RNNs)
 - Motivation
 - Intuition
- Learning with RNNs
 - Formalization
 - Comparison of Feedforward and Recurrent networks
 - Backpropagation through Time (BPTT)
- Problems with RNN Training
 - Vanishing Gradients
 - Exploding Gradients
 - Gradient Clipping

Recurrent Neural Networks



- Up to now
 - Simple neural network structure: 1-to-1 mapping of inputs to outputs
- This lecture: Recurrent Neural Networks
 - Generalize this to arbitrary mappings

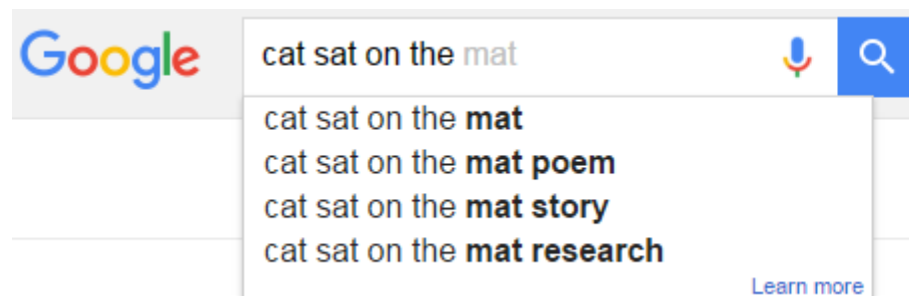

Application: Part-of-Speech Tagging

Legend: Click the legend words to toggle highlighting. [Get help](#) on this page.

Noun Pronoun Verb Adjective Adverb Conjunction Preposition Article Interjection

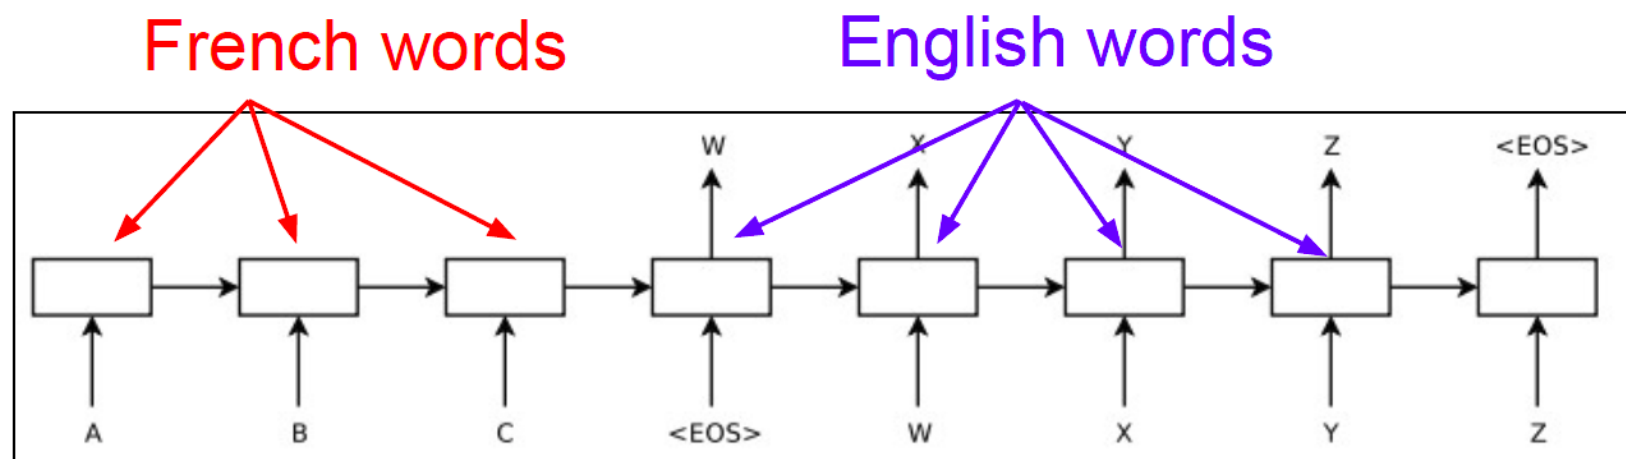
Andrew and Maria thought their jobs were secure after the rancorous argument with the customer , but alas ! Bad news is fast approaching them , especially after they viciously insulted the customer on social media .

Application: Predicting the Next Word



T. Mikolov, M. Karafiat, L. Burget, J. Cernocky, S. Khudanpur, [Recurrent Neural Network Based Language Model](#), Interspeech 2010.

Application: Machine Translation



I. Sutskever, O. Vinyals, Q. Le, [Sequence to Sequence Learning with Neural Networks](#), NIPS 2014.

RNNs: Intuition

- Example: Language modeling

- Suppose we had the training sequence “cat sat on mat”
- We want to train a language model

$$p(\textit{next word} \mid \textit{previous words})$$

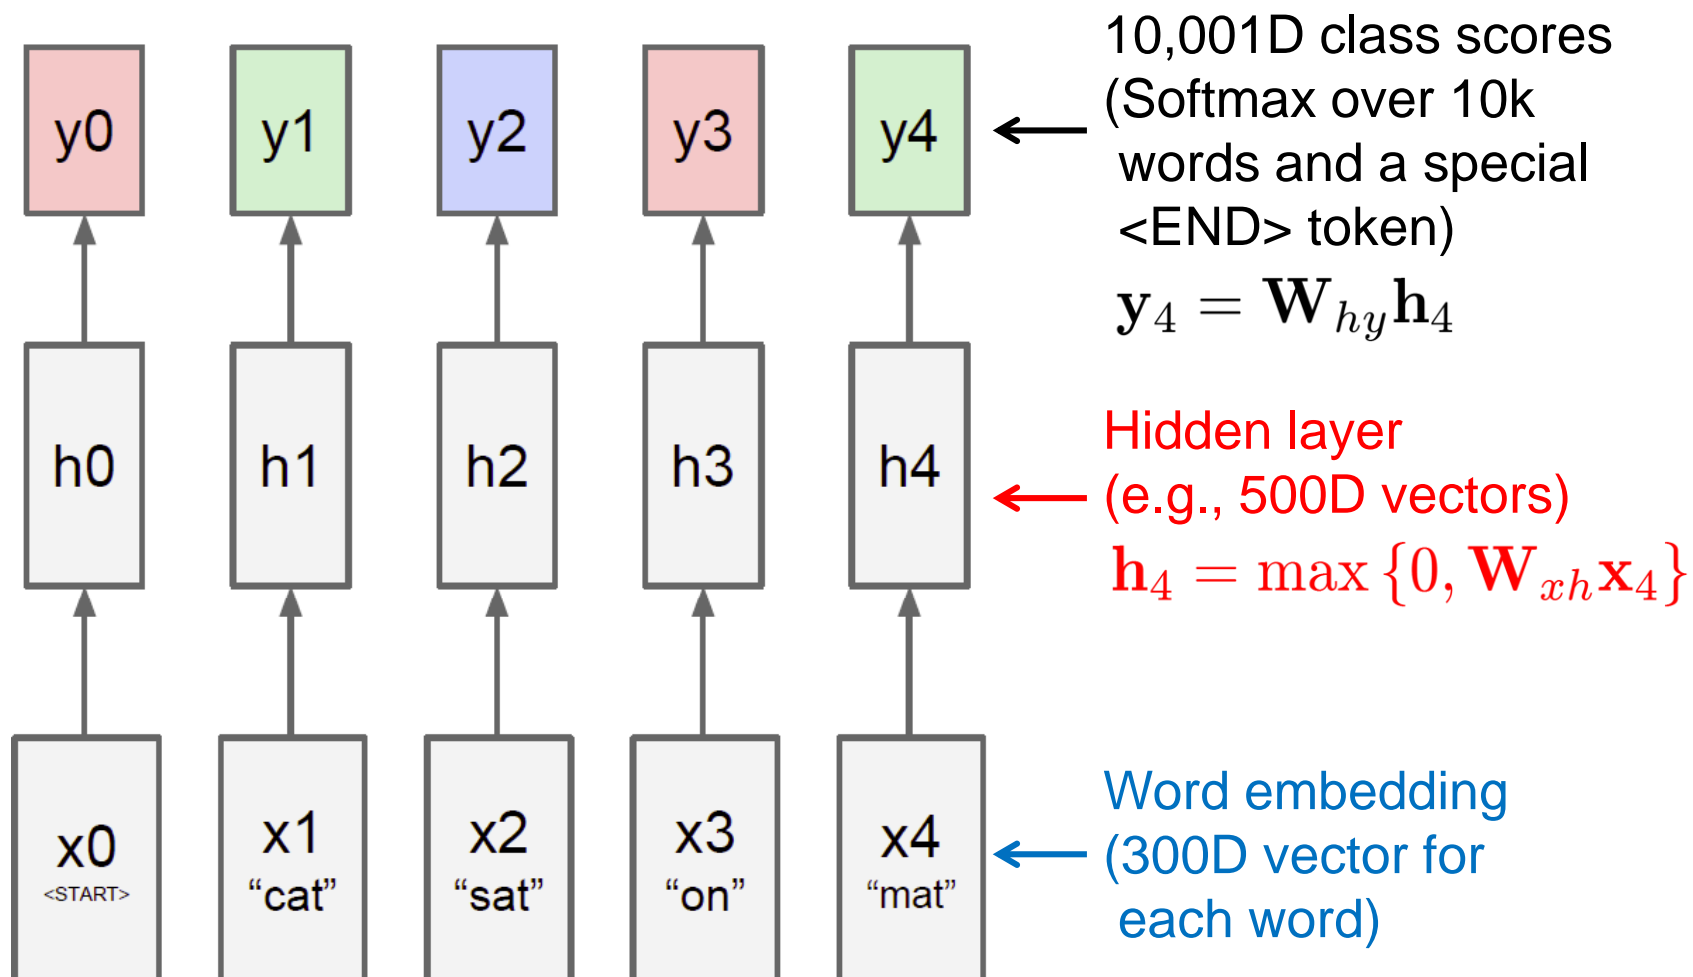
- First assume we only have a finite, 1-word history.
- I.e., we want those probabilities to be high:

- $p(\textit{cat} \mid \langle S \rangle)$
- $p(\textit{sat} \mid \textit{cat})$
- $p(\textit{on} \mid \textit{sat})$
- $p(\textit{mat} \mid \textit{on})$
- $p(\langle E \rangle \mid \textit{mat})$

$\langle S \rangle$ and $\langle E \rangle$ are
start and end tokens.

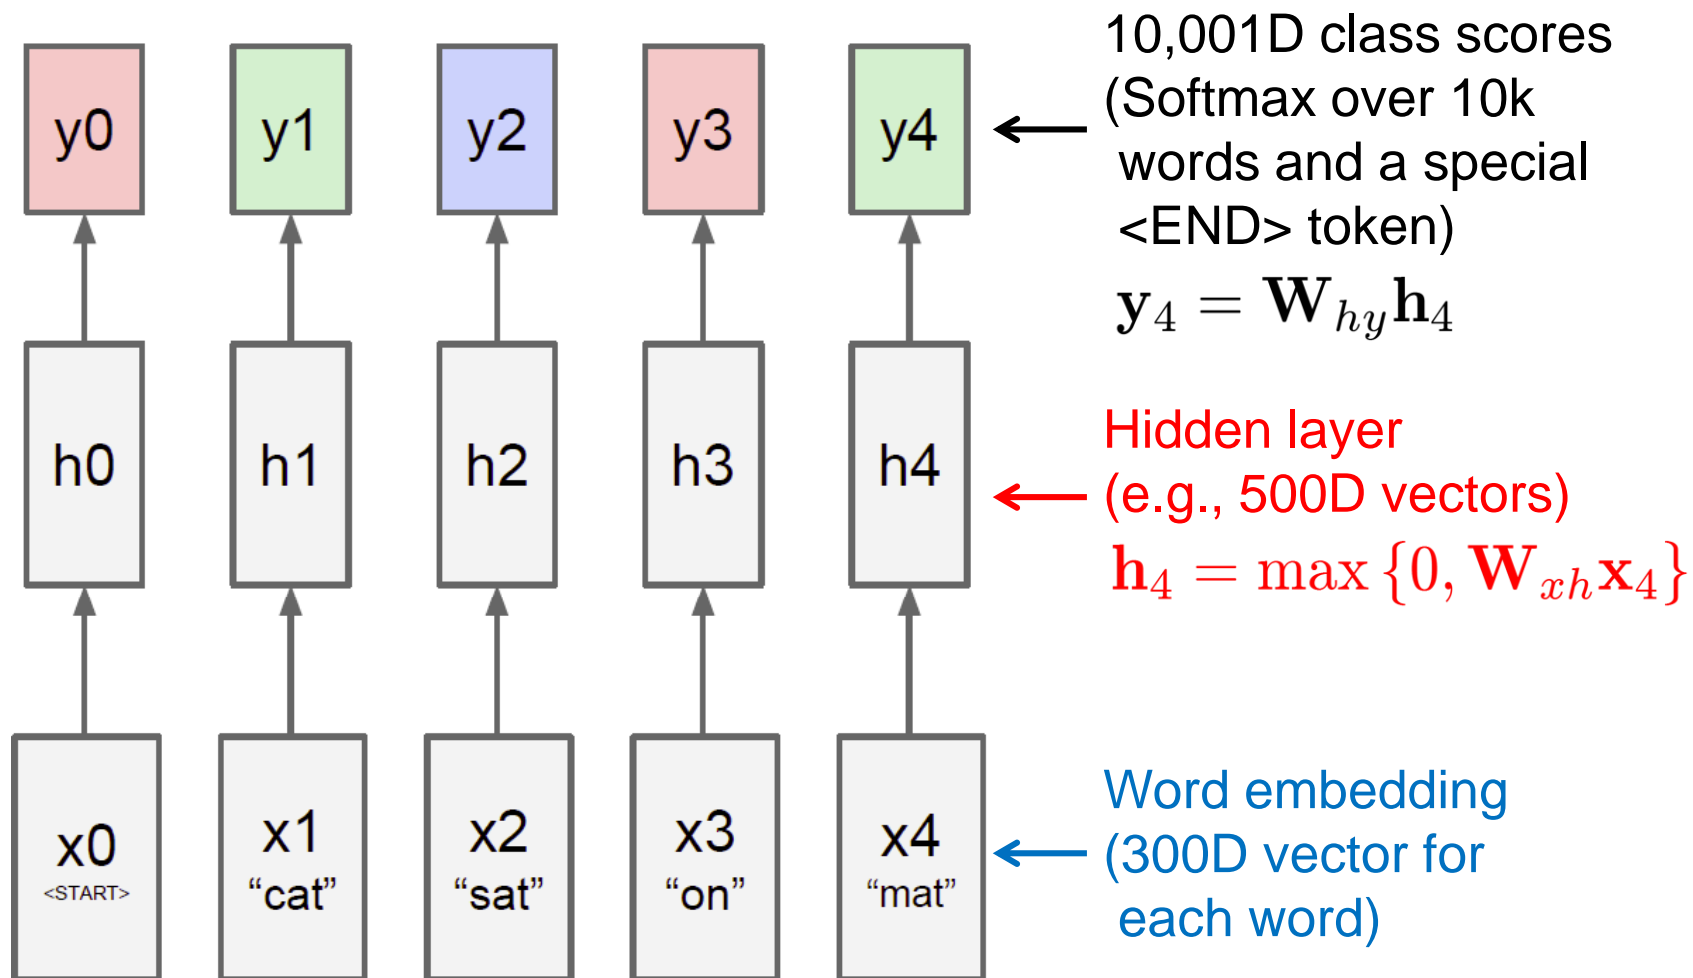
RNNs: Intuition

- Vanilla 2-layer classification net



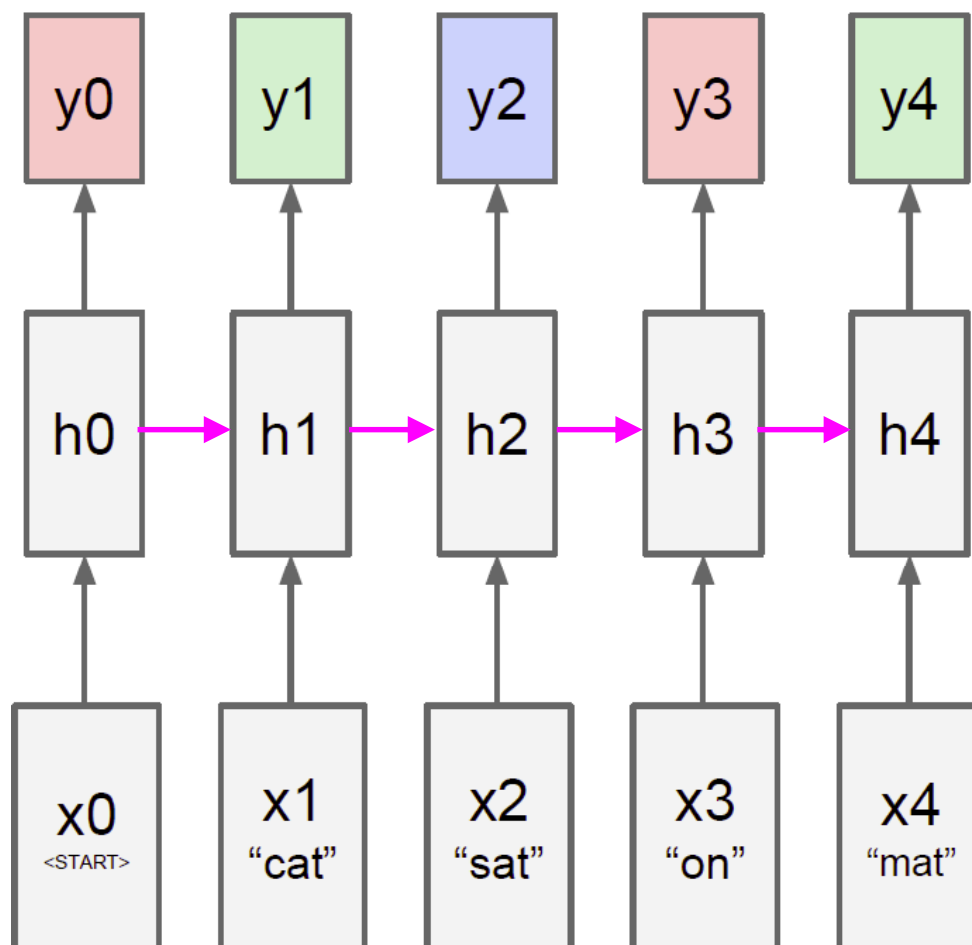
RNNs: Intuition

- Turning this into an RNN (wait for it...)



RNNs: Intuition

- Turning this into an RNN (done!)



10,001D class scores
(Softmax over 10k words and a special <END> token)

$$y_4 = \mathbf{W}_{hy} \mathbf{h}_4$$

Hidden layer
(e.g., 500D vectors)

$$\mathbf{h}_4 = \max \{0, \mathbf{W}_{xh} \mathbf{x}_4 + \mathbf{W}_{hh} \mathbf{h}_3\}$$


Word embedding
(300D vector for each word)

RNNs: Intuition

- Training this on a lot of sentences would give us a language model.

- I.e., a way to predict


$$p(\textit{next word} \mid \textit{previous words})$$



RNNs: Intuition

- Training this on a lot of sentences would give us a language model.
- I.e., a way to predict


$$p(\text{next word} \mid \text{previous words})$$



RNNs: Intuition

- Training this on a lot of sentences would give us a language model.
- I.e., a way to predict

$$p(\textit{next word} \mid \textit{previous words})$$

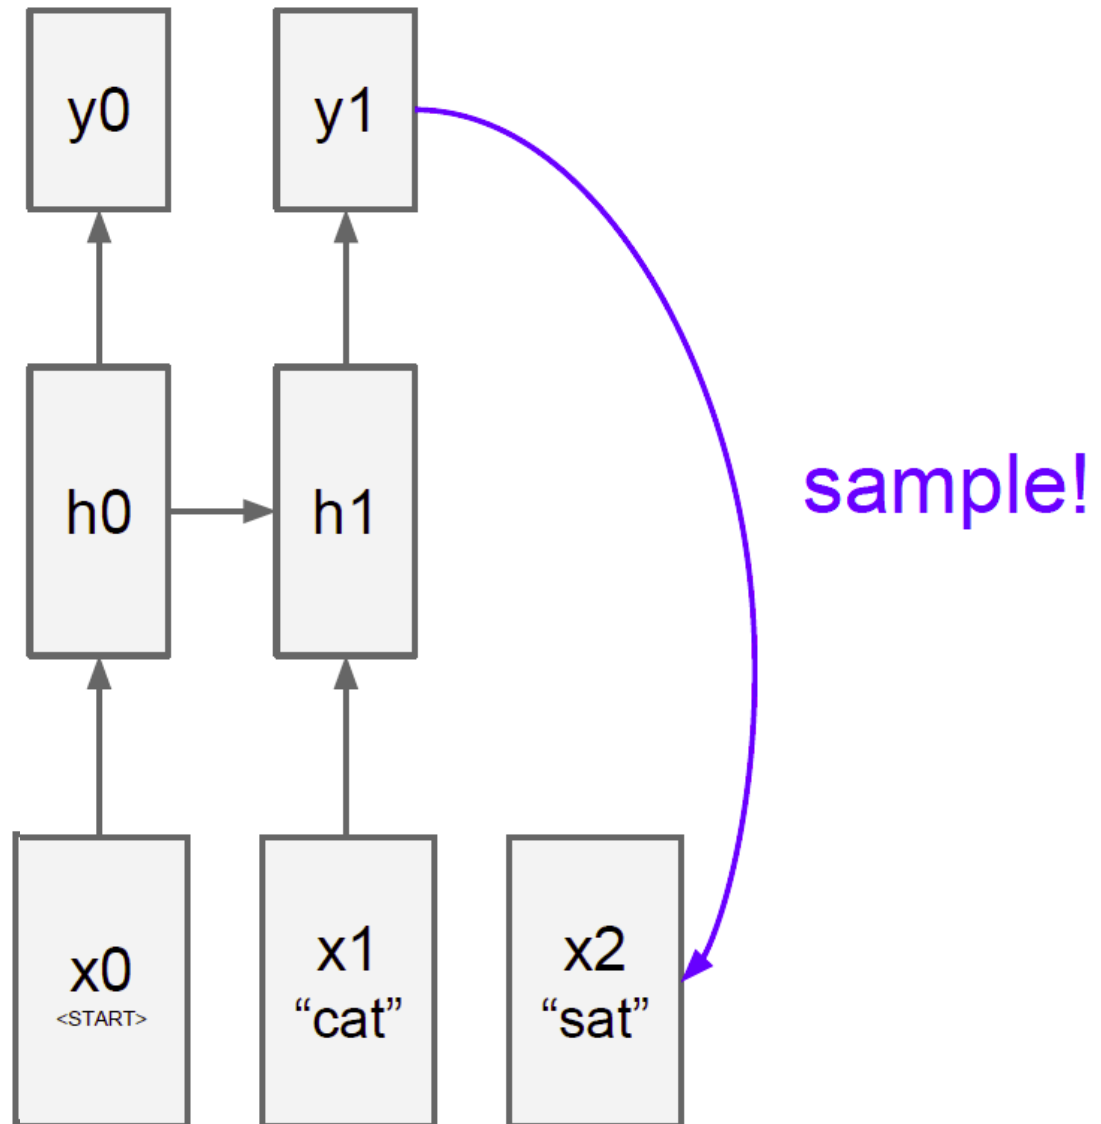


RNNs: Intuition

- Training this on a lot of sentences would give us a language model.

- I.e., a way to predict


$$p(\text{next word} \mid \text{previous words})$$



RNNs: Intuition

- Training this on a lot of sentences would give us a language model.
- I.e., a way to predict

$$p(\text{next word} \mid \text{previous words})$$

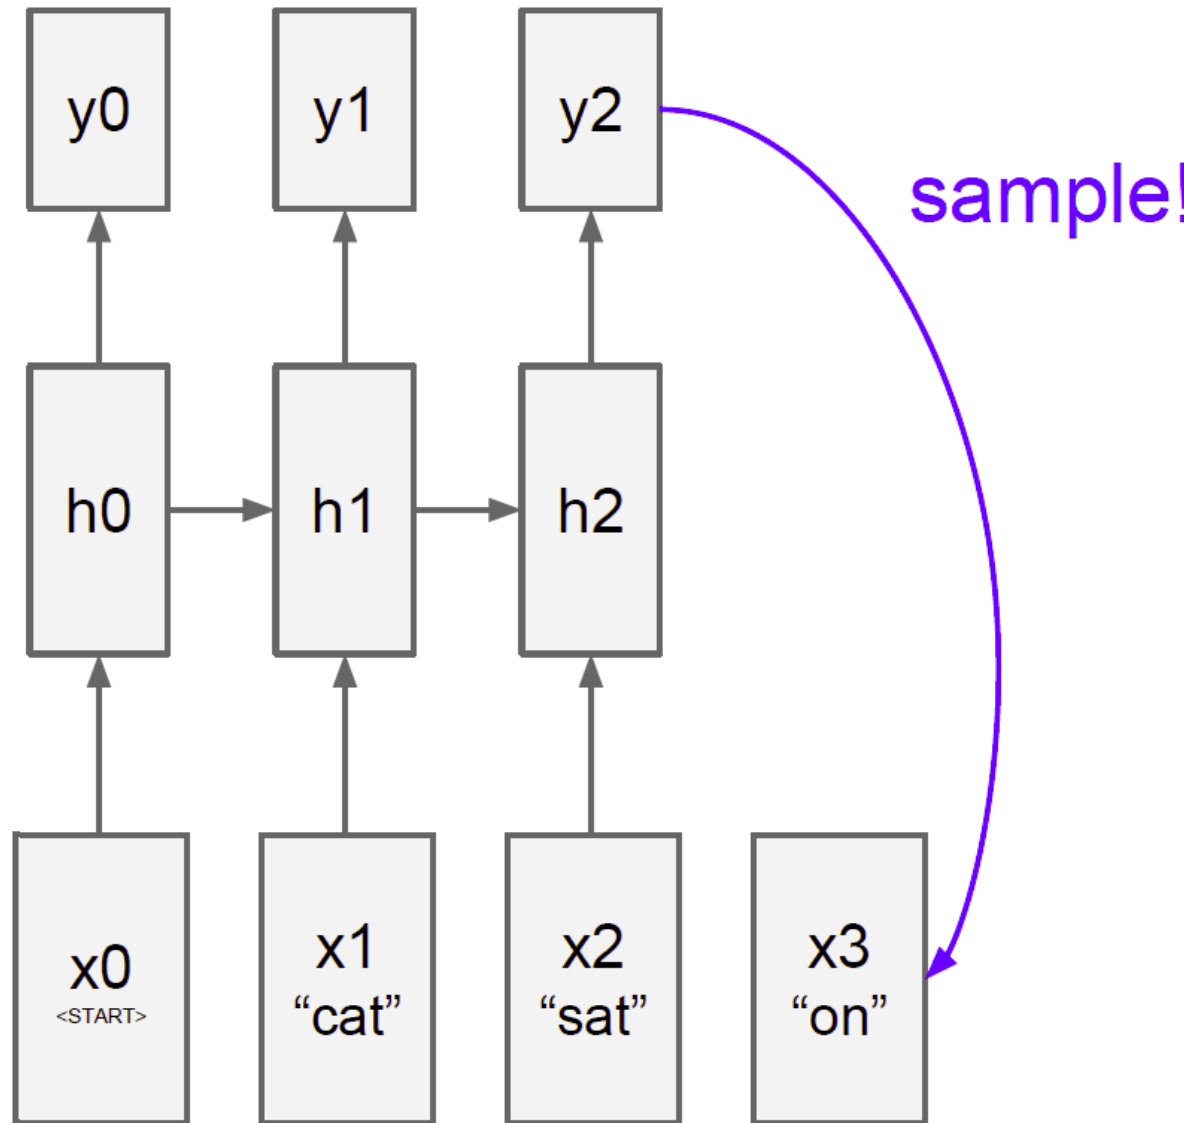


RNNs: Intuition

- Training this on a lot of sentences would give us a language model.


- I.e., a way to predict

$$p(\text{next word} \mid \text{previous words})$$



RNNs: Intuition

- Training this on a lot of sentences would give us a language model.
- I.e., a way to predict
 $p(\text{next word} \mid \text{previous words})$

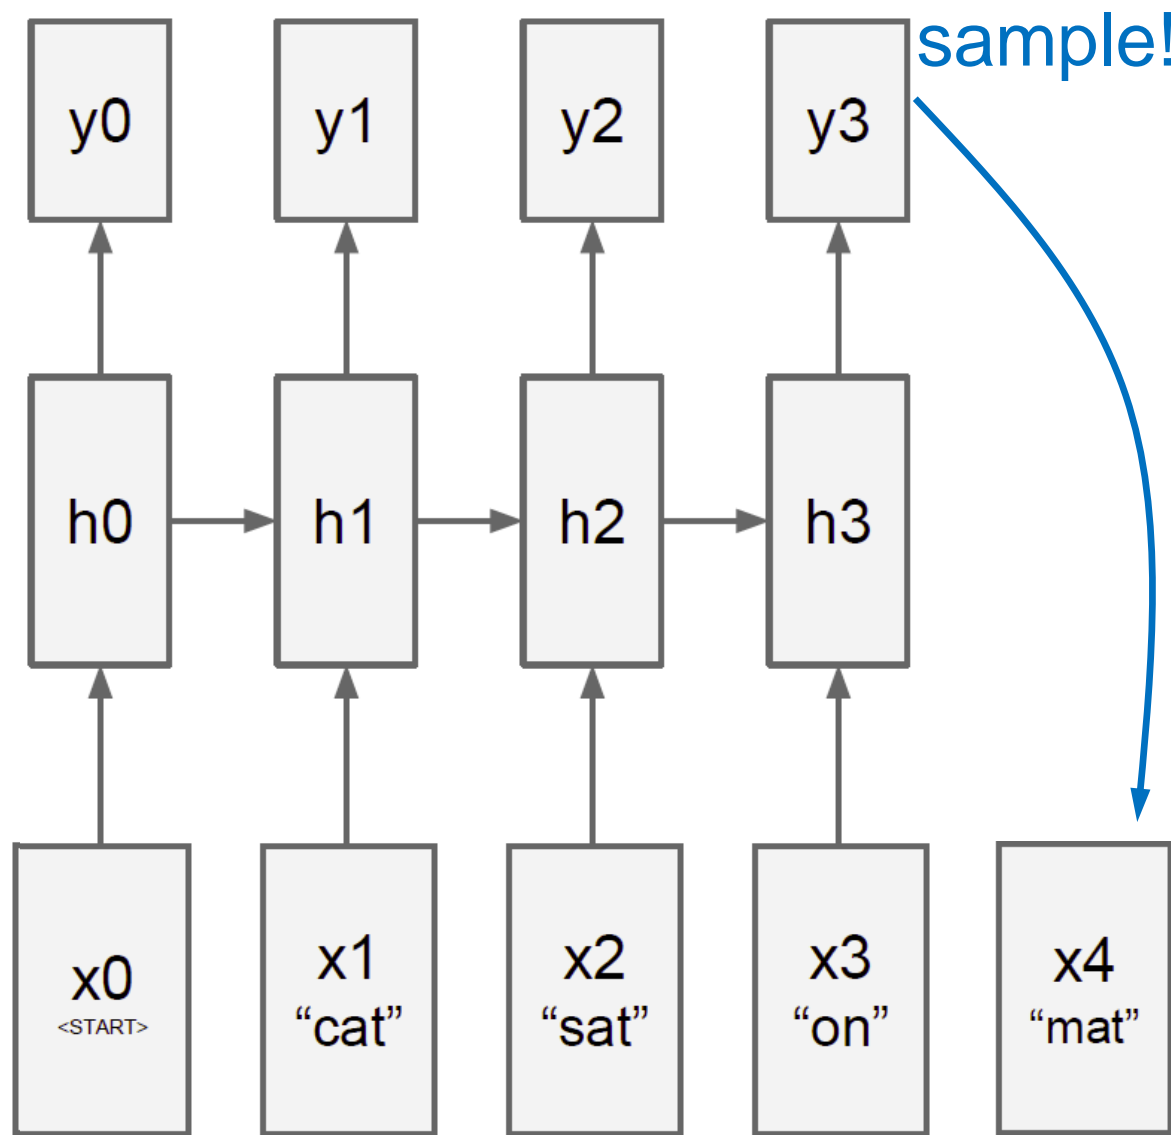


RNNs: Intuition

- Training this on a lot of sentences would give us a language model.

- I.e., a way to predict

$$p(\text{next word} \mid \text{previous words})$$

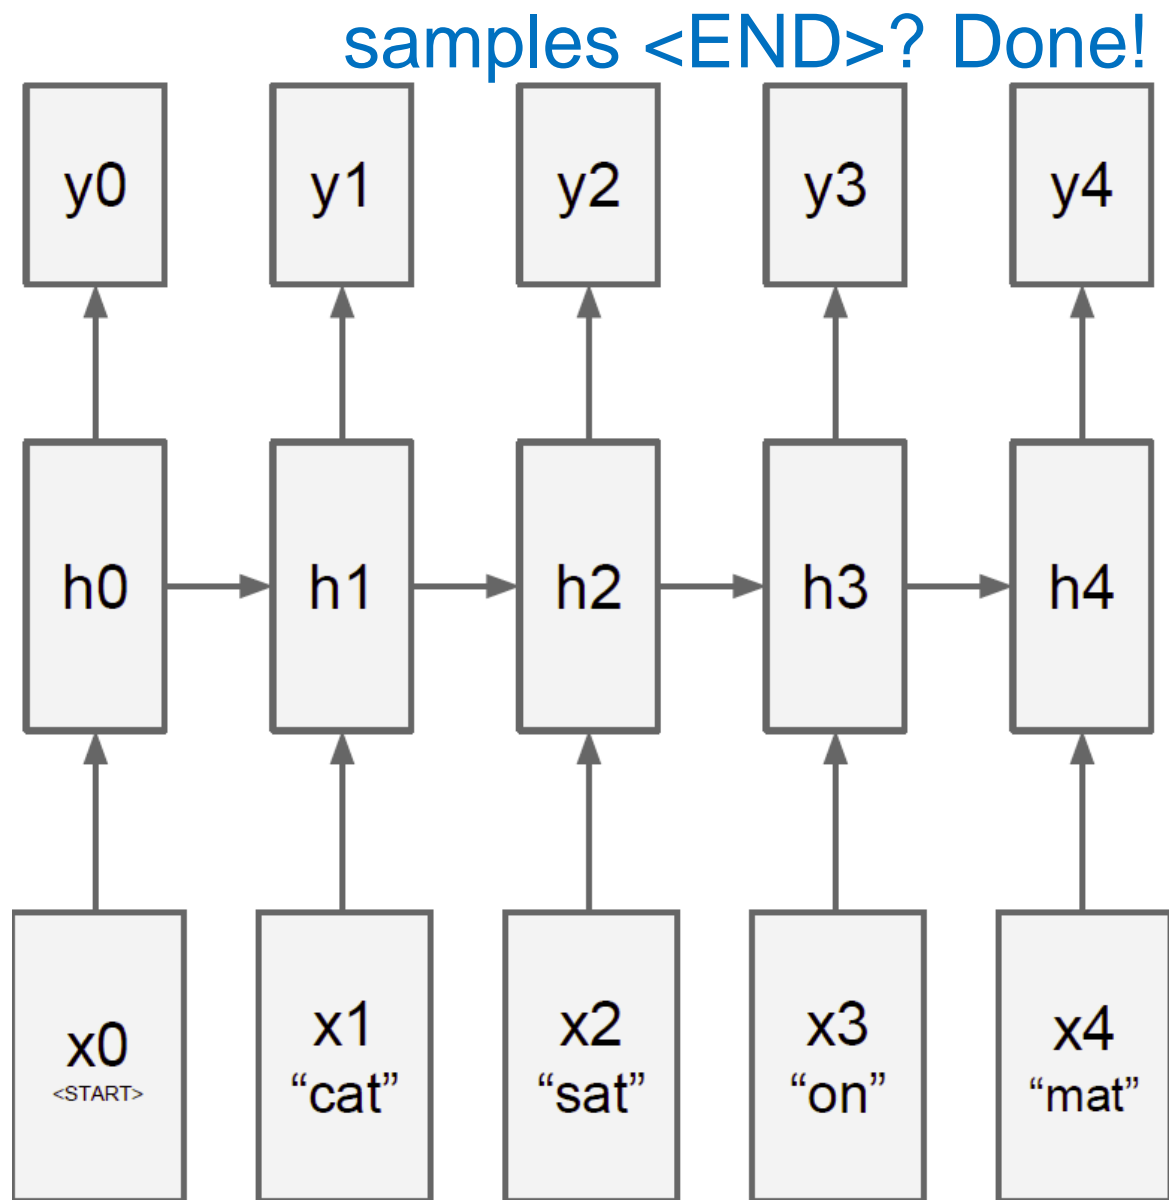


RNNs: Intuition

- Training this on a lot of sentences would give us a language model.

- I.e., a way to predict

$$p(\text{next word} \mid \text{previous words})$$




Topics of This Lecture

- Recurrent Neural Networks (RNNs)
 - Motivation
 - Intuition
- Learning with RNNs
 - Formalization
 - Comparison of Feedforward and Recurrent networks
 - Backpropagation through Time (BPTT)
- Problems with RNN Training
 - Vanishing Gradients
 - Exploding Gradients
 - Gradient Clipping


RNNs: Introduction

- RNNs are regular NNs whose hidden units have additional forward connections over time.
 - You can **unroll** them to create a network that extends over time.
 - When you do this, keep in mind that the weights for the hidden units are shared between temporal layers.



RNNs: Introduction

- RNNs are very powerful, because they combine two properties:
 - Distributed hidden state that allows them to store a lot of information about the past efficiently.
 - Non-linear dynamics that allows them to update their hidden state in complicated ways.
- With enough neurons and time, RNNs can compute anything that can be computed by your computer.




Feedforward Nets vs. Recurrent Nets

- Imagine a feedforward network

- Assume there is a time delay of 1 in using each connection.

⇒ This is very similar to how an RNN works.

- Only change: the layers share their weights.



⇒ The recurrent net is just a feedforward net that keeps reusing the same weights.

Backpropagation with Weight Constraints

- It is easy to modify the backprop algorithm to incorporate linear weight constraints

- To constrain $w_1 = w_2$, we start with the same initialization and then make sure that the gradients are the same:

$$\nabla w_1 = \nabla w_2$$

- We compute the gradients as usual and then use

$$\frac{\partial E}{\partial w_1} + \frac{\partial E}{\partial w_2}$$

for both w_1 and w_2 .

Backpropagation Through Time (BPTT)

• Formalization

- Inputs \mathbf{x}_t
- Outputs \mathbf{y}_t
- Hidden units \mathbf{h}_t
- Initial state \mathbf{h}_0


- Connection matrices

- \mathbf{W}_{xh}
- \mathbf{W}_{hy}
- \mathbf{W}_{hh}


- Configuration

$$\mathbf{h}_t = \sigma(\mathbf{W}_{xh}\mathbf{x}_t + \mathbf{W}_{hh}\mathbf{h}_{t-1} + b)$$

$$\hat{\mathbf{y}}_t = \text{softmax}(\mathbf{W}_{hy}\mathbf{h}_t)$$



Recap: Backpropagation Algorithm




- Efficient propagation scheme

➤ $y_i^{(k-1)}$ is already known from forward pass! (Dynamic Programming)

⇒ Propagate back the gradient from layer k and multiply with $y_i^{(k-1)}$.

Backpropagation Through Time (BPTT)




- Error function

- Computed over all time steps:

$$E = \sum_{1 \leq t \leq T} E_t$$


Backpropagation Through Time (BPTT)



- Backpropagated gradient

➤ For weight w_{ij} :
$$\frac{\partial E_t}{\partial w_{ij}} = \frac{\partial E_t}{\partial \mathbf{h}_t} \frac{\partial \mathbf{h}_t}{\partial w_{ij}}$$

Backpropagation Through Time (BPTT)




- Backpropagated gradient

➤ For weight w_{ij} :

$$\frac{\partial E_t}{\partial w_{ij}} = \frac{\partial E_t}{\partial h_t} \frac{\partial h_t}{\partial w_{ij}} + \frac{\partial E_t}{\partial h_t} \frac{\partial h_t}{\partial h_{t-1}} \frac{\partial h_{t-1}}{\partial w_{ij}}$$

Backpropagation Through Time (BPTT)




- Backpropagated gradient

- For weight w_{ij} :
$$\frac{\partial E_t}{\partial w_{ij}} = \frac{\partial E_t}{\partial \mathbf{h}_t} \frac{\partial \mathbf{h}_t}{\partial w_{ij}} + \frac{\partial E_t}{\partial \mathbf{h}_t} \frac{\partial \mathbf{h}_t}{\partial \mathbf{h}_{t-1}} \frac{\partial \mathbf{h}_{t-1}}{\partial w_{ij}} + \dots$$

- In general:
$$\frac{\partial E_t}{\partial w_{ij}} = \sum_{1 \leq k \leq t} \left(\frac{\partial E_t}{\partial \mathbf{h}_t} \frac{\partial \mathbf{h}_t}{\partial \mathbf{h}_k} \frac{\partial \mathbf{h}_k}{\partial w_{ij}} \right)$$

Backpropagation Through Time (BPTT)




- Analyzing the terms

- For weight w_{ij} :
$$\frac{\partial E_t}{\partial w_{ij}} = \sum_{1 \leq k \leq t} \left(\frac{\partial E_t}{\partial \mathbf{h}_t} \frac{\partial \mathbf{h}_t}{\partial \mathbf{h}_k} \frac{\partial^+ \mathbf{h}_k}{\partial w_{ij}} \right)$$

- This is the “immediate” partial derivative (with \mathbf{h}_{k-1} as constant)

Backpropagation Through Time (BPTT)



- Analyzing the terms

- For weight w_{ij} :

$$\frac{\partial E_t}{\partial w_{ij}} = \sum_{1 \leq k \leq t} \left(\frac{\partial E_t}{\partial \mathbf{h}_t} \frac{\partial \mathbf{h}_t}{\partial \mathbf{h}_k} \frac{\partial^+ \mathbf{h}_k}{\partial w_{ij}} \right)$$

- Propagation term:
- $$\frac{\partial \mathbf{h}_t}{\partial \mathbf{h}_k} = \prod_{t \geq i > k} \frac{\partial \mathbf{h}_i}{\partial \mathbf{h}_{i-1}}$$

Backpropagation Through Time (BPTT)

- Summary

- Backpropagation equations

$$E = \sum_{1 \leq t \leq T} E_t$$

$$\frac{\partial E_t}{\partial w_{ij}} = \sum_{1 \leq k \leq t} \left(\frac{\partial E_t}{\partial h_t} \frac{\partial h_t}{\partial h_k} \frac{\partial^+ h_k}{\partial w_{ij}} \right)$$

$$\frac{\partial h_t}{\partial h_k} = \prod_{t \geq i > k} \frac{\partial \mathbf{h}_i}{\partial \mathbf{h}_{i-1}} = \prod_{t \geq i > k} \mathbf{W}_{hh}^\top \text{diag}(\sigma'(\mathbf{h}_{i-1}))$$

- Remaining issue: how to set the initial state \mathbf{h}_0 ?
- ⇒ Learn this together with all the other parameters.

Topics of This Lecture

- Recurrent Neural Networks (RNNs)
 - Motivation
 - Intuition
- Learning with RNNs
 - Formalization
 - Comparison of Feedforward and Recurrent networks
 - Backpropagation through Time (BPTT)
- Problems with RNN Training
 - Vanishing Gradients
 - Exploding Gradients
 - Gradient Clipping

Problems with RNN Training

- Training RNNs is very hard
 - As we backpropagate through the layers, the magnitude of the gradient may grow or shrink exponentially
⇒ Exploding or vanishing gradient problem!
 - In an RNN trained on long sequences (e.g., 100 time steps) the gradients can easily explode or vanish.
 - Even with good initial weights, it is very hard to detect that the current target output depends on an input from many time-steps ago.

Exploding / Vanishing Gradient Problem

- Consider the propagation equations:

$$\frac{\partial E_t}{\partial w_{ij}} = \sum_{1 \leq k \leq t} \left(\frac{\partial E_t}{\partial h_t} \frac{\partial h_t}{\partial h_k} \frac{\partial^+ h_k}{\partial w_{ij}} \right)$$

$$\begin{aligned} \frac{\partial h_t}{\partial h_k} &= \prod_{t \geq i > k} \frac{\partial \mathbf{h}_i}{\partial \mathbf{h}_{i-1}} = \prod_{t \geq i > k} \mathbf{W}_{hh}^\top \text{diag}(\sigma'(\mathbf{h}_{i-1})) \\ &= (\mathbf{W}_{hh}^\top)^l \end{aligned}$$

- if t goes to infinity and $l = t - k$.

⇒ We are effectively taking the weight matrix to a high power.

- The result will depend on the eigenvalues of \mathbf{W}_{hh} .
 - Largest eigenvalue > 1 ⇒ Gradients *may* explode.
 - Largest eigenvalue < 1 ⇒ Gradients *will* vanish.
 - This is very bad...

Why Is This Bad?

- Vanishing gradients in language modeling
 - Words from time steps far away are not taken into consideration when training to predict the next word.
 - Example:
 - „Jane walked into the room. John walked in too. It was late in the day. Jane said hi to _____“
- ⇒ The RNN will have a hard time learning such long-range dependencies.

Gradient Clipping


- Trick to handle exploding gradients
 - If the gradient is larger than a threshold, clip it to that threshold.

Algorithm 1 Pseudo-code for norm clipping the gradients whenever they explode

$$\begin{aligned} \hat{\mathbf{g}} &\leftarrow \frac{\partial \mathcal{E}}{\partial \theta} \\ \text{if } \|\hat{\mathbf{g}}\| &\geq \textit{threshold} \text{ then} \\ &\quad \hat{\mathbf{g}} \leftarrow \frac{\textit{threshold}}{\|\hat{\mathbf{g}}\|} \hat{\mathbf{g}} \\ \text{end if} \end{aligned}$$

- This makes a big difference in RNNs

Gradient Clipping Intuition



- Example

- Error surface of a single RNN neuron
- High curvature walls
- Solid lines: standard gradient descent trajectories
- Dashed lines: gradients rescaled to fixed size

Handling Vanishing Gradients

- Vanishing Gradients are a harder problem
 - They severely restrict the dependencies the RNN can learn.
 - The problem gets more severe the deeper the network is.
 - It can be very hard to diagnose that Vanishing Gradients occur (you just see that learning gets stuck).
- Ways around the problem
 - Glorot/He initialization (see [Lecture 12](#))
 - ReLU
 - More complex hidden units (LSTM, GRU)

ReLU to the Rescue

- Idea

- Initialize \mathbf{W}_{hh} to identity matrix
- Use Rectified Linear Units (ReLU)



$$g(a) = \max\{0, a\}$$

- Effect

- The gradient is propagated with a constant factor

$$\frac{\partial g(a)}{\partial a} = \begin{cases} 1, & a > 0 \\ 0, & \text{else} \end{cases}$$

⇒ Huge difference in practice!



References and Further Reading

- RNNs
 - R. Pascanu, T. Mikolov, Y. Bengio, [On the difficulty of training recurrent neural networks](#), JMLR, Vol. 28, 2013.
 - A. Karpathy, [The Unreasonable Effectiveness of Recurrent Neural Networks](#), blog post, May 2015.