

Machine Learning – Lecture 20

Recurrent Neural Networks II

15.01.2020

Bastian Leibe


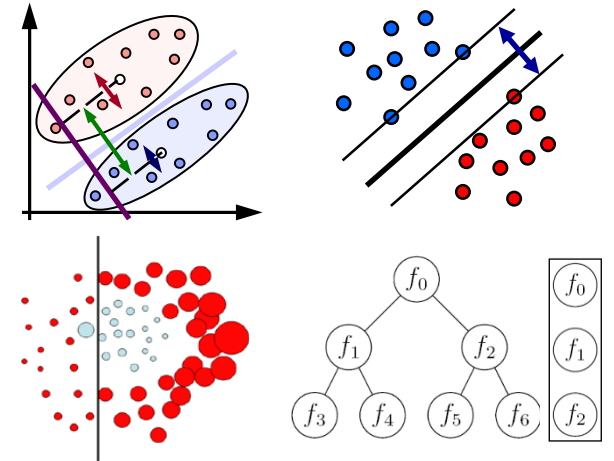

RWTH Aachen

<http://www.vision.rwth-aachen.de>

leibe@vision.rwth-aachen.de

Course Outline

- Fundamentals
 - Bayes Decision Theory
 - Probability Density Estimation
- Classification Approaches
 - Linear Discriminants
 - Support Vector Machines
 - Ensemble Methods & Boosting
 - Random Forests
- Deep Learning
 - Foundations
 - Convolutional Neural Networks
 - Recurrent Neural Networks




Topics of This Lecture


- **Recap: Recurrent Neural Networks (RNNs)**
 - Backpropagation through Time (BPTT)
 - Problems with RNN Training
 - Handling Vanishing Gradients
- Improved hidden units for RNNs
 - Long Short-Term Memory (LSTM)
 - Gated Recurrent Units (GRU)
- Applications of RNNs

Recurrent Neural Networks


one to one




one to many




many to one



many to many



many to many




- Up to now
 - Simple neural network structure: 1-to-1 mapping of inputs to outputs
- This lecture: Recurrent Neural Networks
 - Generalize this to arbitrary mappings

Recap: Recurrent Neural Networks (RNNs)


- RNNs are regular NNs whose hidden units have additional connections over time.

- You can **unroll** them to create a network that extends over time.
- When you do this, keep in mind that the weights for the hidden are shared between temporal layers.



- RNNs are very powerful
 - With enough neurons and time, they can compute anything that can be computed by your computer.

Recap: Backpropagation Through Time (BPTT)



- Configuration
- Backpropagated gradient


$$\mathbf{h}_t = \sigma(\mathbf{W}_{xh}\mathbf{x}_t + \mathbf{W}_{hh}\mathbf{h}_{t-1} + b)$$

$$\hat{\mathbf{y}}_t = \text{softmax}(\mathbf{W}_{hy}\mathbf{h}_t)$$

- For weight w_{ij} :

$$\frac{\partial E_t}{\partial w_{ij}} = \sum_{1 \leq k \leq t} \left(\frac{\partial E_t}{\partial \mathbf{h}_t} \frac{\partial \mathbf{h}_t}{\partial \mathbf{h}_k} \frac{\partial^+ \mathbf{h}_k}{\partial w_{ij}} \right)$$

Recap: Backpropagation Through Time (BPTT)




- Analyzing the terms

- For weight w_{ij} :
$$\frac{\partial E_t}{\partial w_{ij}} = \sum_{1 \leq k \leq t} \left(\frac{\partial E_t}{\partial h_t} \frac{\partial h_t}{\partial h_k} \frac{\partial^+ h_k}{\partial w_{ij}} \right)$$

- This is the “immediate” partial derivative (with \mathbf{h}_{k-1} as constant)

Recap: Backpropagation Through Time (BPTT)



- Analyzing the terms

- For weight w_{ij} :

$$\frac{\partial E_t}{\partial w_{ij}} = \sum_{1 \leq k \leq t} \left(\frac{\partial E_t}{\partial h_t} \frac{\partial h_t}{\partial h_k} \frac{\partial^+ h_k}{\partial w_{ij}} \right)$$

- Propagation term:
- $$\frac{\partial h_t}{\partial h_k} = \prod_{t \geq i > k} \frac{\partial h_i}{\partial h_{i-1}}$$

Recap: Exploding / Vanishing Gradient Problem

- BPTT equations:

$$\frac{\partial E_t}{\partial w_{ij}} = \sum_{1 \leq k \leq t} \left(\frac{\partial E_t}{\partial h_t} \frac{\partial h_t}{\partial h_k} \frac{\partial^+ h_k}{\partial w_{ij}} \right)$$

$$\begin{aligned} \frac{\partial h_t}{\partial h_k} &= \prod_{t \geq i > k} \frac{\partial \mathbf{h}_i}{\partial \mathbf{h}_{i-1}} = \prod_{t \geq i > k} \mathbf{W}_{hh}^\top \text{diag}(\sigma'(\mathbf{h}_{i-1})) \\ &= (\mathbf{W}_{hh}^\top)^l \end{aligned}$$

(if t goes to infinity and $l = t - k$.)

⇒ We are effectively taking the weight matrix to a high power.

- The result will depend on the eigenvalues of \mathbf{W}_{hh} .
 - Largest eigenvalue $> 1 \Rightarrow$ Gradients *may* explode.
 - Largest eigenvalue $< 1 \Rightarrow$ Gradients *will* vanish.
 - This is very bad...

Understanding Exploding/Vanishing Gradients

- What changes when going from $\mathbf{W}_{hh}^\top \frac{\partial E}{\partial h_t}$ to $(\mathbf{W}_{hh}^\top)^k \frac{\partial E}{\partial h_t}$?

- Eigendecomposition of the matrix

$$\mathbf{W}_{hh}^\top = \mathbf{Q} \begin{bmatrix} \lambda_1 & & 0 \\ & \ddots & \\ 0 & & \lambda_d \end{bmatrix} \mathbf{Q}^{-1} = \mathbf{Q} \mathbf{\Lambda} \mathbf{Q}^{-1}$$

$$(\mathbf{W}_{hh}^\top)^2 = \mathbf{Q} \mathbf{\Lambda} \mathbf{Q}^{-1} \mathbf{Q} \mathbf{\Lambda} \mathbf{Q}^{-1} = \mathbf{Q} \mathbf{\Lambda}^2 \mathbf{Q}^{-1}$$

$$\vdots$$

$$(\mathbf{W}_{hh}^\top)^k = \mathbf{Q} \mathbf{\Lambda}^k \mathbf{Q}^{-1}$$

⇒ Effect: The eigenvalues are taken to a high power.

- This is bad, because

- If $\lambda_i < 1$, then $\lambda_i^k \rightarrow 0$
- If $\lambda_i > 1$, then $\lambda_i^k \rightarrow \infty$

$(\mathbf{W}_{hh}^\top)^k$ either attenuates
or amplifies the gradient
by an exponential factor!

Why Is This Bad?

- Vanishing gradients in language modeling
 - Words from time steps far away are not taken into consideration when training to predict the next word.
 - Example:
 - „Jane walked into the room. John walked in too. It was late in the day. Jane said hi to _____“
- ⇒ The RNN will have a hard time learning such long-range dependencies.

Gradient Clipping


- Trick to handle exploding gradients
 - If the gradient is larger than a threshold, clip it to that threshold.

Algorithm 1 Pseudo-code for norm clipping the gradients whenever they explode

$$\begin{aligned} \hat{\mathbf{g}} &\leftarrow \frac{\partial \mathcal{E}}{\partial \theta} \\ \text{if } \|\hat{\mathbf{g}}\| &\geq \textit{threshold} \text{ then} \\ &\quad \hat{\mathbf{g}} \leftarrow \frac{\textit{threshold}}{\|\hat{\mathbf{g}}\|} \hat{\mathbf{g}} \\ \text{end if} \end{aligned}$$

- This makes a big difference in RNNs

Gradient Clipping Intuition



- Example

- Error surface of a single RNN neuron
- High curvature walls
- Solid lines: standard gradient descent trajectories
- Dashed lines: gradients rescaled to fixed size

Handling Vanishing Gradients

- Vanishing Gradients are a harder problem
 - They severely restrict the dependencies the RNN can learn.
 - The problem gets more severe the deeper the network is.
 - It can be very hard to diagnose that Vanishing Gradients occur (you just see that learning gets stuck).
- Ways around the problem
 - Glorot/He initialization (see [Lecture 15](#))
 - ReLU
 - More complex hidden units ([LSTM](#), [GRU](#))

ReLU to the Rescue

- Idea

- Initialize \mathbf{W}_{hh} to identity matrix
- Use Rectified Linear Units (ReLU)



$$g(a) = \max\{0, a\}$$

- Effect

- The gradient is propagated with a constant factor

$$\frac{\partial g(a)}{\partial a} = \begin{cases} 1, & a > 0 \\ 0, & \text{else} \end{cases}$$

⇒ Huge difference in practice!



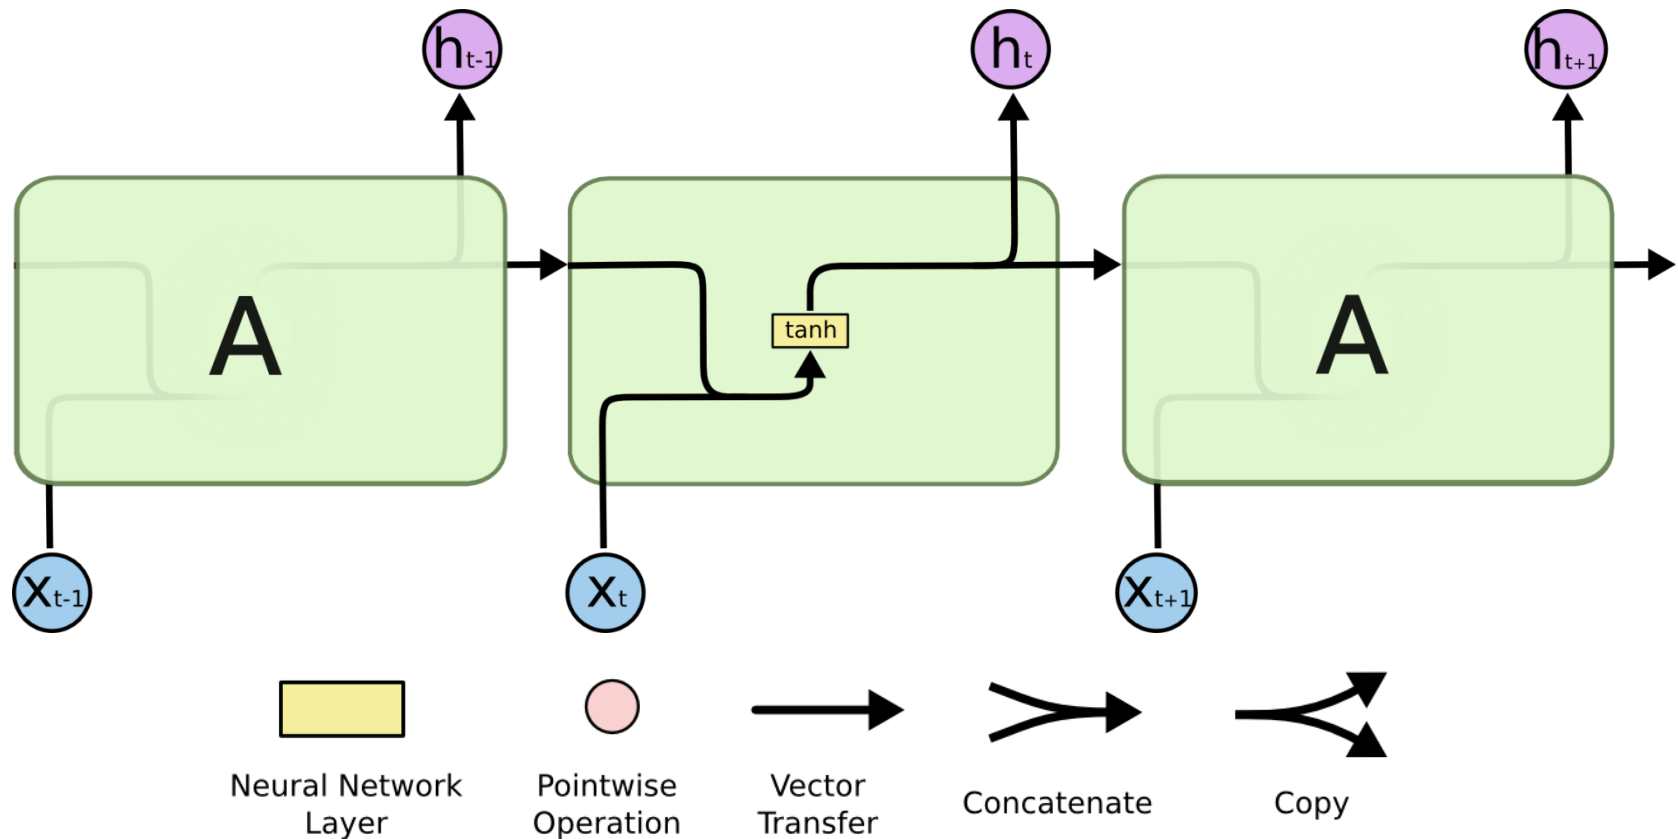
Topics of This Lecture

- Recap: Recurrent Neural Networks (RNNs)
 - Backpropagation through Time (BPTT)
 - Problems with RNN Training
 - Handling Vanishing Gradients
- Improved hidden units for RNNs
 - Long Short-Term Memory (LSTM)
 - Gated Recurrent Units (GRU)
- Applications of RNNs

More Complex Hidden Units

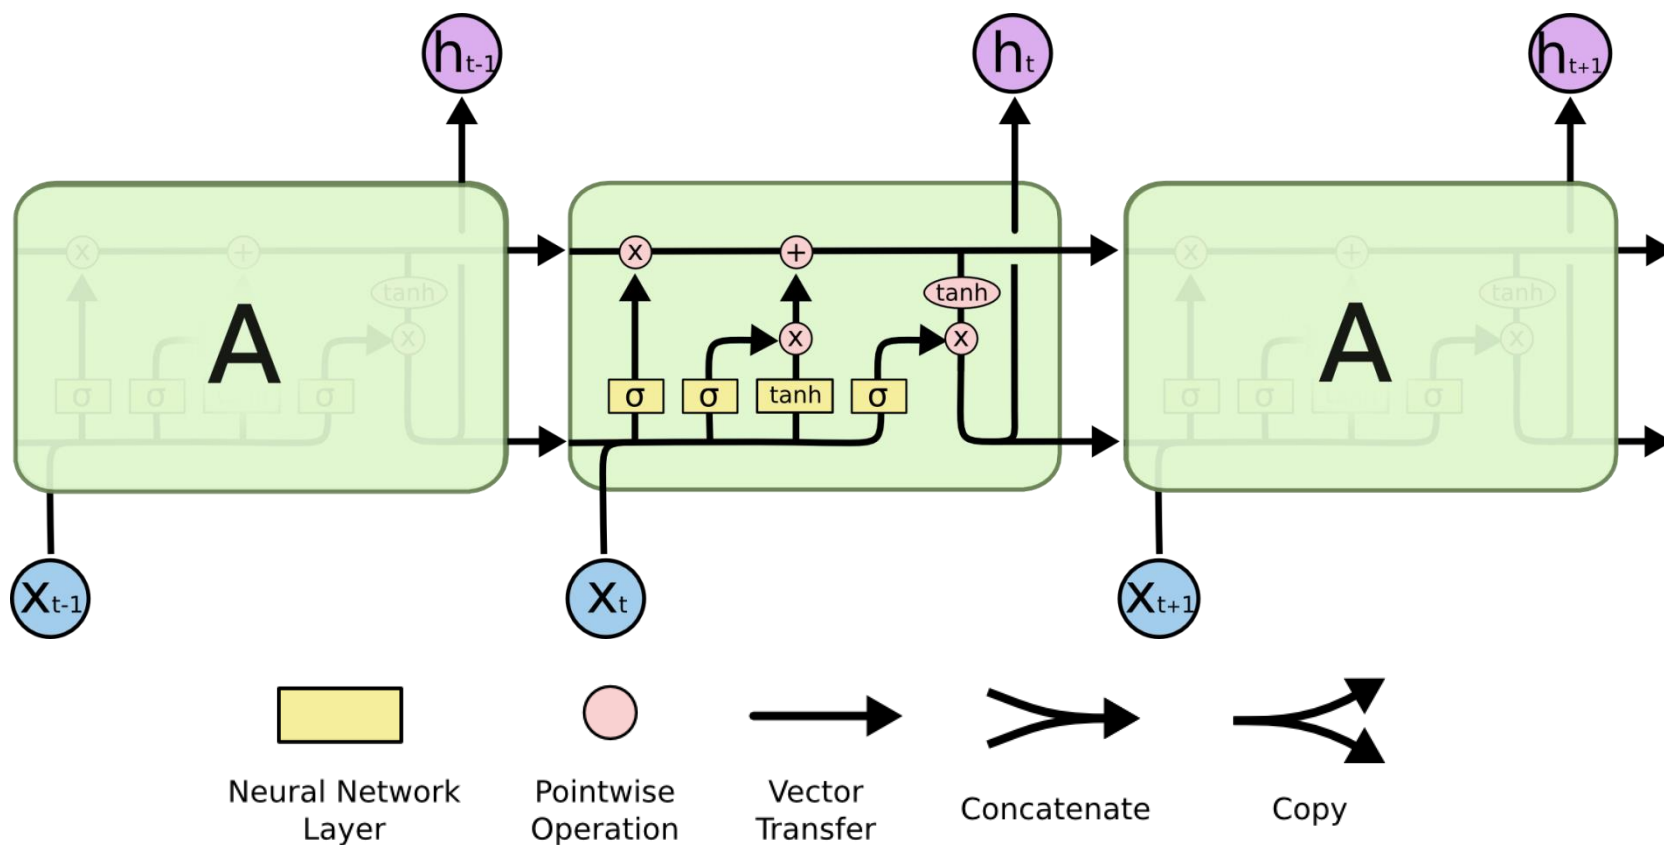
- Target properties
 - Want to achieve constant error flow through a single unit
 - At the same time, want the unit to be able to pick up long-term connections or focus on short-term ones, as the problem demands.
- Ideas behind LSTMs
 - Take inspiration from the design of memory cells
 - Keep around memories to capture long distance dependencies
 - Allow error messages to flow at different strengths depending on the inputs

Memory Representation in RNNs



- RNNs can be seen as chains of repeating modules
 - In a standard RNN, the repeating module has a very simple structure (e.g., a tanh)

Long Short-Term Memory




- LSTMs

- Repeating modules have 4 layers, interacting in a special way.

LSTMs: Core Ideas


- Cell state

- This is the key to LSTMs.
- It acts like a conveyor belt, information can flow along it unchanged.



- Gates


- The cell state can be modified through gates.
- Structure: sigmoid net layer + pointwise multiplication
- The sigmoid outputs values between 0 and 1
 - 0: Let nothing through
 - 1: Let everything through
- The gate layers are learned together with all other parameters.



Elements of LSTMs

- Forget gate layer

- Look at \mathbf{h}_{t-1} and \mathbf{x}_t and output a number between 0 and 1 for each dimension in the cell state \mathbf{C}_{t-1} .
0: completely delete this,
1: completely keep this.



$$f_t = \sigma (W_f \cdot [h_{t-1}, x_t] + b_f)$$


- Example

- Task: try to predict the next word
 - Cell state could include the gender of the present subject
- ⇒ When we see a new subject, want to forget the gender of the old subject.

Elements of LSTMs

- **Update gate layer**

- Decide what information to store in the cell state.
- Sigmoid network (**input gate layer**) decides which values are updated.
- tanh layer creates a vector of new candidate values \tilde{C}_t that could be added to the state.



$$i_t = \sigma(W_i \cdot [h_{t-1}, x_t] + b_i)$$


$$\tilde{C}_t = \tanh(W_C \cdot [h_{t-1}, x_t] + b_C)$$

- **In the example**

- Add the gender of the new subject to the cell state.

Elements of LSTMs

- Updating the state
 - Multiply the old state by f_t , forgetting the things we decided to forget.
 - Then add $i_t * \tilde{C}_t$ the new candidate values, scaled by how much we decided to update each value.




$$C_t = f_t * C_{t-1} + i_t * \tilde{C}_t$$

- In the example
 - Combined effect: replace the old gender by the new one.

Elements of LSTMs

- **Output gate layer**

- Output is a filtered version of our gate state.
- First, apply sigmoid layer to decide what parts of the cell state to output.
- Then, pass the cell state through a tanh (to push the values to be between -1 and 1) and multiply it with the output of the sigmoid gate.



$$o_t = \sigma (W_o [h_{t-1}, x_t] + b_o)$$


$$h_t = o_t * \tanh (C_t)$$

- **In the example**

- Since we just saw a subject, might want to output information relevant to a verb (e.g., whether the subject is singular or plural).


RNN vs. LSTM

- LSTM just changes the form of the equation for h such that:
 1. More expressive multiplicative interactions become possible
 2. Gradients flow nicer
 3. The network can explicitly decide to reset the hidden state
- Those changes have a huge effect in practice
 - LSTMs perform much better than regular RNNs
 - Many applications have become possible with LSTMs that weren't feasible before.




LSTMs in Practice

- LSTMs are currently highly en vogue
 - Popular default model for most sequence labeling tasks.
 - Very powerful, especially when stacked and made even deeper.
 - Most useful if you have lots and lots of data.
- ⇒ Very active research field
- Here are also some other ways of illustrating them




1/14/20



Extension: Gated Recurrent Units (GRU)

- Simpler model than LSTM
 - Combines the forget and input gates into a single **update gate** z_t .
 - Similar definition for a **reset gate** r_t , but with different weights.
 - In both cases, merge the cell state and hidden state.



$$z_t = \sigma(W_z \cdot [h_{t-1}, x_t])$$

$$r_t = \sigma(W_r \cdot [h_{t-1}, x_t])$$

$$\tilde{h}_t = \tanh(W \cdot [r_t * h_{t-1}, x_t])$$

$$h_t = (1 - z_t) * h_{t-1} + z_t * \tilde{h}_t$$

- Empirical results
 - Performance similar to LSTM (no clear winner yet)
 - But GRU has fewer parameters.

GRUs: Intuition

- Effects

- If reset is close to 0, ignore previous hidden state.
⇒ Allows model to drop information that is irrelevant in the future.
- Update gate z controls how much of past state should matter now.
⇒ If z is close to 0, then we can copy information in that unit through many time steps!
⇒ Less vanishing gradients!


$$z_t = \sigma(W_z \cdot [h_{t-1}, x_t])$$

$$r_t = \sigma(W_r \cdot [h_{t-1}, x_t])$$

$$\tilde{h}_t = \tanh(W \cdot [r_t * h_{t-1}, x_t])$$

$$h_t = (1 - z_t) * h_{t-1} + z_t * \tilde{h}_t$$

GRUs: Intuition



$$z_t = \sigma(W_z \cdot [h_{t-1}, x_t])$$

$$r_t = \sigma(W_r \cdot [h_{t-1}, x_t])$$

$$\tilde{h}_t = \tanh(W \cdot [r_t * h_{t-1}, x_t])$$

$$h_t = (1 - z_t) * h_{t-1} + z_t * \tilde{h}_t$$

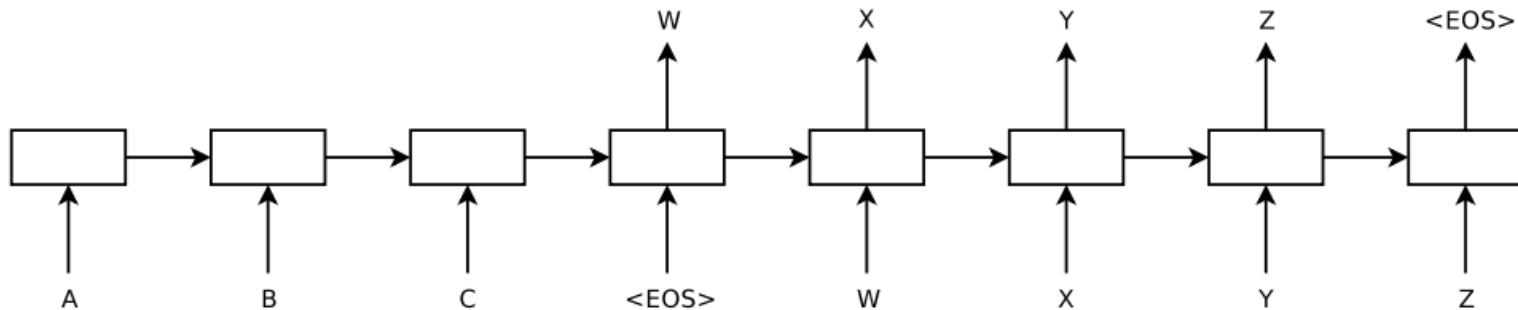
- Typical learned behaviors
 - Units with short-term dependencies often have active reset gate
 - Units with long-term dependencies have inactive update gates.

Topics of This Lecture

- Recap: Recurrent Neural Networks (RNNs)
 - Backpropagation through Time (BPTT)
 - Problems with RNN Training
 - Handling Vanishing Gradients
- Improved hidden units for RNNs
 - Long Short-Term Memory (LSTM)
 - Gated Recurrent Units (GRU)
- Applications of RNNs

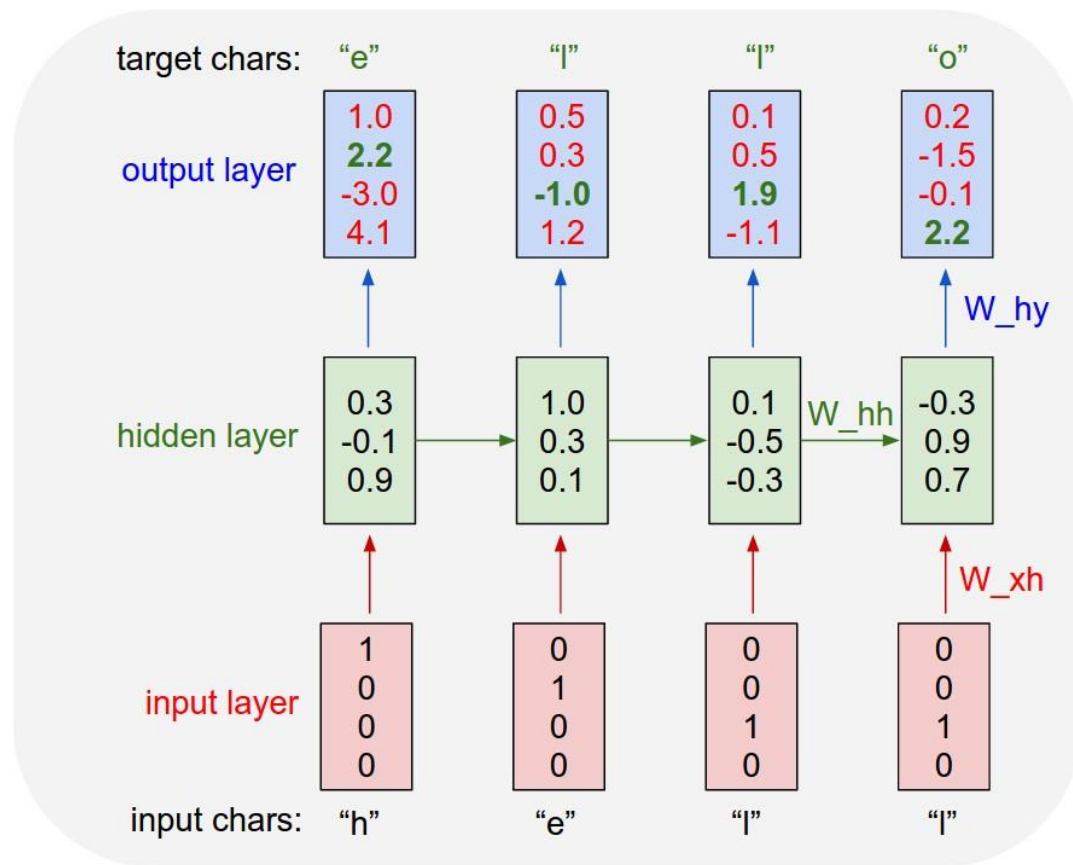
Applications

- Machine Translation [Sutskever et al., 2014]



Application: Character-Level Language Model

- Setup
 - RNN trained on huge amounts of text
 - Task: model the prob. distribution of the next character in the sequence.
- Main advantage of RNN here
 - RNN can learn varying amount of context



Language Model Results

PANDARUS:

Alas, I think he shall be come approached and the day
When little strain would be attain'd into being never fed,
And who is but a chain and subjects of his death,
I should not sleep.

Second Senator:

They are away this miseries, produced upon my soul,
Breaking and strongly should be buried, when I perish
The earth and thoughts of many states.

DUKE VINCENTIO:

Well, your wit is in the care of side and that.

Second Lord:

They would be ruled after this chamber, and
my fair nudes begun out of the fact, to be conveyed,
Whose noble souls I'll have the heart of the wars.

- Example: Generating Shakespeare
 - Trained on all works of Shakespeare (4.4 MB of data)
 - Using a 3-Layer RNN with 512 hidden units per layer

Language Model Results

Naturalism and decision for the majority of Arab countries' capitalide was grounded by the Irish language by [[John Clair]], [[An Imperial Japanese Revolt]], associated with Guangzham's sovereignty. His generals were the powerful ruler of the Portugal in the [[Protestant Immineners]], which could be said to be directly in Cantonese Communication, which followed a ceremony and set inspired prison, training. The emperor travelled back to [[Antioch, Perth, October 25|21]] to note, the Kingdom of Costa Rica, unsuccessful fashioned the [[Thrales]], [[Cynth's Dajoard]], known in western [[Scotland]], near Italy to the conquest of India with the conflict. Copyright was the succession of independence in the slop of Syrian influence that was a famous German movement based on a more popular servicious, non-doctrinal and sexual power post. Many governments recognize the military housing of the [[Civil Liberalization and Infantry Resolution 265 National Party in Hungary]], that is sympathetic to be to the [[Punjab Resolution]] (PJS)[<http://www.humah.yahoo.com/guardian.cfm/7754800786d17551963s89.htm> Official economics Adjoint for the Nazism, Montgomery was swear to advance to the resources for those Socialism's rule, was starting to signing a major tripad of aid exile.]]

- Example: Generating Wikipedia pages
 - Trained on 100MB of Wikipedia data
 - Using an LSTM

Language Model Results

For $\bigoplus_{n=1,\dots,m}$ where $\mathcal{L}_{m\bullet} = 0$, hence we can find a closed subset \mathcal{H} in \mathcal{H} and any sets \mathcal{F} on X , U is a closed immersion of S , then $U \rightarrow T$ is a separated algebraic space.

Proof. Proof of (1). It also start we get

$$S = \mathrm{Spec}(R) = U \times_X U \times_X U$$

and the comparicoly in the fibre product covering we have to prove the lemma generated by $\coprod Z \times_U U \rightarrow V$. Consider the maps M along the set of points Sch_{fppf} and $U \rightarrow U$ is the fibre category of S in U in Section, ?? and the fact that any U affine, see Morphisms, Lemma ??. Hence we obtain a scheme S and any open subset $W \subset U$ in $Sh(G)$ such that $\mathrm{Spec}(R') \rightarrow S$ is smooth or an

$$U = \bigcup U_i \times_{S_i} U_i$$

which has a nonzero morphism we may assume that f_i is of finite presentation over S . We claim that $\mathcal{O}_{X,x}$ is a scheme where $x, x', s'' \in S'$ such that $\mathcal{O}_{X,x'} \rightarrow \mathcal{O}'_{X',x'}$ is


- Example: Hallucinating Algebraic Geometry
 - Trained on an Algebraic Geometry book
 - Using a multilayer LSTM

Language Model Results

```
/*
 * Increment the size file of the new incorrect UI_FILTER group information
 * of the size generatively.
 */
static int indicate_policy(void)
{
    int error;
    if (fd == MARN_EPT) {
        /*
         * The kernel blank will coeld it to userspace.
         */
        if (ss->segment < mem_total)
            unblock_graph_and_set_blocked();
        else
            ret = 1;
        goto bail;
    }
    segaddr = in_SB(in.addr);
    selector = seg / 16;
    setup_works = true;
    for (i = 0; i < blocks; i++) {
        seq = buf[i++];
        bpf = bd->bd.next + i * search;
        if (fd) {
```

- Example:
Hallucinating C Code
 - Trained on the Linux source code (474MB from github)
 - Using a large 3-layer LSTM

Applications: Image Tagging



- Simple combination of CNN and RNN
 - Use CNN to define initial state \mathbf{h}_0 of an RNN.
 - Use RNN to produce text description of the image.

Applications: Image Tagging

- Setup
 - Train on corpus of images with textual descriptions
 - E.g. Microsoft CoCo
 - 120k images
 - 5 sentences each

a man riding a bike on a dirt path through a forest.
bicyclist raises his fist as he rides on desert dirt trail.
this dirt bike rider is smiling and raising his fist in triumph.
a man riding a bicycle while pumping his fist in the air.
a mountain biker pumps his fist in celebration.



Results: Image Tagging



a group of people standing
around a room with
remotes
logprob: -9.17



a young boy is holding a
baseball bat
logprob: -7.61



a cow is standing in the middle of a street
logprob: -8.84

Spectacular results!

Results: Image Tagging



a baby laying on a bed with a stuffed bear
logprob: -8.66




a young boy is holding a
baseball bat
logprob: -7.65



a cat is sitting on a couch with a remote control
logprob: -12.45

- Wrong, but one can still see why those results were selected...

Results: Image Tagging



a toilet with a seat up in a bathroom
logprob: -13.44



a woman holding a teddy bear in front of a mirror
logprob: -9.65



a horse is standing in the middle of a road
logprob: -10.34

- Not sure what happened here...

Fun Application: Image to Story



Later on the eighth day , Billy was a friend of a man who lived on his own . He did n't know how far away they were , and if he was to survive the fall . His mind raced , trying not to show any signs of weakness . The wind ruffled the snow and ice in the snow . He had no idea how many times he was going to climb into the mountains . He told me to stay on the ground for a while , but if I find out what s going on , we should go on foot . Sam and Si Lei joined us in the army .

- Example: Generating a story from an image
 - Trained on corpus of adventure novels

More Results




Having lain on the bed , I did n't know what to say . He turned his attention to the room and saw a large room . The room was furnished with a single bed , a dresser and a large bed with a table in the center of the room . It was a long time ago . The room was designed with the most powerful and efficient ones . As far as I m concerned , it was a long time ago . On the other side of the room was a beautiful picture of a woman who had been abducted by the fireplace and their own personal belongings in order to keep it safe , but it didn t take too long . Feeling helpless , he turned his attention back to me . ``

More Results



Only Prince Darin knew how to run from the mountains , and once more , he could see the outline of a rider on horseback . The wind ruffled his hair in an attempt to locate the forest . He hadn't been in such a state of mind before , but it was a good thing . All of them seemed to be doing the same thing . They didn't know where they came from . The wind blew up the mountain peaks and disappeared into the sky , leaving trails behind the peaks of the mountains on Mount Fuji .

Application: Video to Text Description



Video-to-Text Results

Correct descriptions.



S2VT: A man is doing stunts on his bike.



S2VT: A herd of zebras are walking in a field.



S2VT: A young woman is doing her hair.



S2VT: A man is shooting a gun at a target.

Relevant but incorrect descriptions.



S2VT: A small bus is running into a building.



S2VT: A man is cutting a piece of a pair of a paper.



S2VT: A cat is trying to get a small board.



S2VT: A man is spreading butter on a tortilla.

Irrelevant descriptions.



S2VT: A man is pouring liquid in a pan.



S2VT: A polar bear is walking on a hill.



S2VT: A man is doing a pencil.



S2VT: A black clip to walking through a path.

References and Further Reading

- RNNs

- R. Pascanu, T. Mikolov, Y. Bengio, [On the difficulty of training recurrent neural networks](#), JMLR, Vol. 28, 2013.
- A. Karpathy, [The Unreasonable Effectiveness of Recurrent Neural Networks](#), blog post, May 2015.

- LSTM

- S. Hochreiter, J. Schmidhuber, [Long short-term memory](#), Neural Computation, Vol. 9(8): 1735–1780, 1997.
- A. Graves, [Generating Sequences With Recurrent Neural Networks](#), ArXiv 1308.0850v5, 2014.
- C. Olah, [Understanding LSTM Networks](#), blog post, August 2015.